



Introduction to R

Ali Zaidi, Machine Learning and Data Science Education Team

May 17, 2016

Overview of the R Project

Lesson Plan

R U Ready?

- What the R Language for Statistical Computing is
- R's capabilities and its limitations
- What types of problems you might want to use R with
- How to manage data with open source R
- How to develop models and write functions in R



What is R?

Why should I care?

- R is a successor of the S Language, originated at Bell Labs AT&T
- Based on the Scheme interpreter
- Originally designed by two University of Auckland Professors for their introduction to statistics course



R Philosophy

What R Thou?

R follows the [Unix philosophy](#)

Write programs that do one thing and do it well. Write programs to work together.

- R is extensible with more than 9,000 packages available at CRAN (<http://crantastic.org/packages>)
- R, like its inspiration, Scheme, is a functional programming language
- R is lazy, and lazy evaluation can be used to interface to other languages
- R is a highly interpreted dynamically typed language, allowing you to mutate variables and analyze datasets quickly, but is significantly slower than low-level, statically typed languages like C or Java
- R has a high memory footprint, and can easily lead to crashes if you aren't careful



Development Environments

Where to Write R Code

- The most popular integrated development environment for R is [RStudio](#)
- The RStudio IDE is entirely html/javascript based, so completely cross-platform
- RStudio Server for cloud instances
- Developers of RStudio have also written a plethora of useful R packages
- For Windows machines, we have recently announced the general availability of [R Tools for Visual Studio, RTVS](#)
- RTVS will support connectivity to Azure and SQL Server very soon
- RTVS has great debugging support



Quick Tour of Your IDE

Strengths of R

Where R Succeeds

- Expressive
- Open source
- Extendable – nearly 1000 packages with functions to use
- Focused on statistics and machine learning – utilized by academics and practitioners
- Advanced data structures and graphical capabilities
- Large user community, academics and industry
- It is designed by statisticians



Weaknesses of R

Where R Falls Short

- It is designed by statisticians
- Inefficient at element-by-element computations
- May make large demands on system resources, namely memory
- Data capacity limited by memory
- Single-threaded



Some Essential Open Source Packages

- There are over 1,000 R packages to choose from, what do I start with?
- Data Management: `dplyr`, `tidyr`, `data.table`
- Visualization: `ggplot2`, `ggvis`, `htmlwidgets`, `shiny`
- Data Importing: `haven`, `RODBC`
- Other favorites; `magrittr`, `rmarkdown`, `caret`



R Foundations

Command line prompts

Symbol	Meaning
>	ready for a new command
+	awaiting the completion of an existing command



I'm Lost!

Getting Help with R

- [Stack Overflow](#)
- [Cross Validated, R](#)
- [R Reference Card](#)
- [RStudio Cheat Sheets](#)
- [R help mailing list and archives](#)
- [CRAN Task Views](#)
- [Crantastic](#)
- [Revolutions Blog](#)
- [R-Bloggers](#)



Quick Tour of Things You Need to Know

Data Structures

- R's data structures can be described by their dimensionality, and their type.

	Homogeneous	Heterogeneous
1d	Atomic vector	List
2d	Matrix	Data frame
nd	Array	



Quick Tour of Things You Need to Know

Data Types

- Atomic vectors come in one of four types
- `logical` (boolean). Values: `TRUE` | `FALSE`
- `integer`
- `double` (often called numeric)
- `character`
- Rare types:
- `complex`
- `raw`



Lab 1: R Data Types

Data Manipulation with the dplyr Package

Overview

At the end of this session, you will have learned:

- How to manipulate data quickly with **dplyr** using a very intuitive 'grammar'
- How to use **dplyr** to perform common exploratory and manipulation procedures
- Apply your own custom functions to group manipulations **dplyr** with `mutate()`, `summarise()` and `do()`
- Connect to remote databases to work with larger than memory datasets



Why use dplyr?

The Grammar of Data Manipulation

- R comes with a plethora of base functions for data manipulation
- **dplyr** makes data manipulation easier by providing a few functions for the most common tasks and procedures
- **dplyr** achieves remarkable speed-up gains by using a C++ backend
- **dplyr** has multiple backends for working with data stored in various sources: SQLite, MySQL, bigquery, and more
- **dplyr** was inspired to give data manipulation a simple, cohesive grammar (similar philosophy to **ggplot** - grammar of graphics)
- the recent package **dplyrXdf** brings much of the same functionality of **dplyr** to XDF data



Manipulation verbs

filter

select rows based on matching criteria

slice

select rows by number

select

select columns by column names

arrange

reorder rows by column values

mutate

add new variables based on transformations of existing variables

transmute

transform and drop other variables



Aggregation verbs

group_by

identify grouping variables for calculating groupwise summary statistics

count

count the number of records per group

summarise

calculate one or more summary functions per group, returning one row of results per group (or one for the entire dataset)



Viewing Data

- `dplyr` includes a wrapper called `tbl_df` makes `df` into a 'local `df`' that improves the printing of dataframes in the console
- if you want to see more of the data you can still coerce to `data.frame`

```
library(dplyr)
bankData <- read.table("data/bank-full.csv",
                      header = T, sep = ";",
                      stringsAsFactors = F)
(bankData <- tbl_df(bankData))
```

```
## Source: local data frame [45,211 x 17]
##
##   age      job  marital education default balance housing loan
## 1   58 management married  tertiary      no    2143      yes   no
## 2   44 technician single secondary      no      29      yes   no
## 3   33 entrepreneur married secondary      no       2      yes  yes
## 4   47 blue-collar married  unknown      no   1506      yes   no
## 5   33      unknown single   unknown      no       1       no   no
## 6   35 management married  tertiary      no    231      yes   no
## 7   28 management single   tertiary      no   447      yes  yes
## 8   42 entrepreneur divorced tertiary     yes       2      yes   no
## 9   58      retired married   primary      no    121      yes   no
## 10  43 technician single secondary      no    593      yes   no
## .. ...      ...      ...      ...      ...      ...      ...
## Variables not shown: contact (chr), day (int), month (chr), duration
## (int), campaign (int), pdays (int), previous (int), poutcome (chr), y
## (chr)
```



Filtering and Reordering Data

Subsetting Data

- `dplyr` makes subsetting by rows very easy
- The `filter` verb takes conditions for filtering rows based on conditions

```
filter(bankData,  
  month %in% c("april", "may", "jun"),  
  default == "yes")
```

```
## Source: local data frame [351 x 17]
```

```
##
```

```
##   age      job   marital education default balance housing loan  
## 1  42 entrepreneur divorced  tertiary      yes         2     yes   no  
## 2  55    services divorced secondary      yes         1     yes   no  
## 3  51     admin.   single secondary      yes        -2     no    no  
## 4  33 technician married secondary      yes        72     yes   no  
## 5  33 blue-collar single secondary      yes       -60     no    no  
## 6  60    retired married secondary      yes        15     no    no  
## 7  35 entrepreneur married secondary      yes       204     yes   no  
## 8  41 blue-collar single primary       yes      -137     yes  yes  
## 9  54    services married secondary      yes         0     yes   no  
## 10 52 blue-collar divorced primary       yes      -183     yes   no  
## .. ...      ...      ...      ...      ...      ...      ... ..
```

```
## Variables not shown: contact (chr), day (int), month (chr), duration
```

```
## (int), campaign (int), pdays (int), previous (int), poutcome (chr), y
```

```
## (chr)
```



Exercise

Your turn:

- How many defaults in the dataset?
- How many of the entrepreneurs that defaulted were also divorced?



Solution

```
defaults <- filter(bankData, default == "yes")  
nrow(defaults)
```

```
## [1] 815
```

```
brokeEntrepreneurs <- filter(bankData,  
                             job == "entrepreneur",  
                             marital == "divorced",  
                             default == "yes")  
  
nrow(brokeEntrepreneurs)
```

```
## [1] 6
```



Select a set of columns

- You can use the `select()` verb to specify which columns of a dataset you want
- This is similar to the `keep` option in SAS's data step.
- Use a colon `:` to select all the columns between two variables (inclusive)
- Use `contains` to take any columns containing a certain word/phrase/character



Select Example 1

```
select(bankData, age, job, default, balance, housing)
```

```
## Source: local data frame [45,211 x 5]
```

```
##
```

```
##   age      job default balance housing
## 1  58 management    no   2143    yes
## 2  44 technician    no    29    yes
## 3  33 entrepreneur  no     2    yes
## 4  47 blue-collar    no  1506    yes
## 5  33      unknown    no     1     no
## 6  35 management    no   231    yes
## 7  28 management    no   447    yes
## 8  42 entrepreneur  yes     2    yes
## 9  58      retired    no   121    yes
## 10 43 technician    no   593    yes
## .. ...           ...    ...    ...
```



Select: Other Options

starts_with(x, ignore.case = FALSE)

name starts with x

ends_with(x, ignore.case = FALSE)

name ends with x

matches(x, ignore.case = FALSE)

selects all variables whose name matches the regular expression x

num_range("V", 1:5, width = 1)

selects all variables (numerically) from **v1** to **v5**.

You can also use a - to drop variables.



Reordering Data

- You can reorder your dataset based on conditions using the `arrange()` verb

```
arrange(bankData, balance, default)
```

```
## Source: local data frame [45,211 x 17]
##
##   age      job      marital education default balance housing loan
## 1  26 blue-collar single  secondary     yes   -8019      no   yes
## 2  49  management married  tertiary     yes   -6847      no   yes
## 3  60  management divorced tertiary     no   -4057     yes   no
## 4  43  management married  tertiary     yes   -3372     yes   no
## 5  57 self-employed married  tertiary     yes   -3313     yes   yes
## 6  39 self-employed married  tertiary     no   -3058     yes   yes
## 7  40  technician married  tertiary     yes   -2827     yes   yes
## 8  52  management married  tertiary     no   -2712     yes   yes
## 9  49 blue-collar single   primary     yes   -2604     yes   no
## 10 51  management divorced tertiary     no   -2282     yes   yes
## .. ...
## Variables not shown: contact (chr), day (int), month (chr), duration
##   (int), campaign (int), pdays (int), previous (int), poutcome (chr), y
##   (chr)
```



Exercise

Use `arrange()` to sort on the basis of `y`, `marital`, `job` (in descending order), and `balance`



Solution

```
arrange(bankData, y, marital, desc(job), balance)
```

```
## Source: local data frame [45,211 x 17]
```

```
##
```

```
##   age    job  marital education default balance housing loan  contact
## 1  26 unknown divorced secondary    no   -295    yes   no  cellular
## 2  47 unknown divorced  primary    no     0     no   no  cellular
## 3  82 unknown divorced   unknown    no     0     no   no  telephone
## 4  59 unknown divorced   unknown    no    27     no   no   unknown
## 5  31 unknown divorced tertiary    no   137     no   no  cellular
## 6  56 unknown divorced  primary    no   558     no   no   unknown
## 7  51 unknown divorced   unknown    no  1649     no   no   unknown
## 8  66 unknown divorced   unknown    no  1993    yes   no  cellular
## 9  56 unknown divorced   unknown    no  2152     no   no   unknown
## 10 53 unknown divorced tertiary    no  2272    yes   no   unknown
## .. ...      ...      ...      ...      ...      ...      ...      ...
```

```
## Variables not shown: day (int), month (chr), duration (int), campaign
```

```
##   (int), pdays (int), previous (int), poutcome (chr), y (chr)
```



Summary

filter

Extract subsets of rows. See also `slice()`

select

Extract subsets of columns. See also `rename()`

arrange

Sort your data



Transformations

- The `mutate()` verb can be used to make new columns

```
mutate(bankData,  
      DefaultFlag = ifelse(default == 'yes',  
                            yes = 1,  
                            no = 0)  
)
```

```
## Source: local data frame [45,211 x 18]  
##  
##   age      job   marital education default balance housing loan  
## 1   58 management married  tertiary      no    2143     yes   no  
## 2   44 technician single   secondary      no     29     yes   no  
## 3   33 entrepreneur married  secondary      no      2     yes  yes  
## 4   47 blue-collar married   unknown      no   1506     yes   no  
## 5   33      unknown single    unknown      no      1      no   no  
## 6   35 management married  tertiary      no    231     yes   no  
## 7   28 management single   tertiary      no    447     yes  yes  
## 8   42 entrepreneur divorced tertiary     yes      2     yes   no  
## 9   58      retired married   primary      no    121     yes   no  
## 10  43 technician single   secondary      no    593     yes   no  
## .. ...  
## Variables not shown: contact (chr), day (int), month (chr), duration  
##   (int), campaign (int), pdays (int), previous (int), poutcome (chr), y  
##   (chr), DefaultFlag (dbl)
```



Summarise Data by Groups

- The `group_by` verb creates a grouping by a categorical variable
- Functions can be placed inside `summarise` to create summary functions

```
summarise(group_by(bankData, default), Num = n())
```

```
## Source: local data frame [2 x 2]
```

```
##
```

```
##   default   Num
```

```
## 1      no 44396
```

```
## 2      yes  815
```



Example group_by 2

```
summarise(group_by(bankData, default), ave_balance = mean(balance))
```

```
## Source: local data frame [2 x 2]  
##  
##   default ave_balance  
## 1      no   1389.8064  
## 2     yes   -137.6245
```



Example group_by 3

```
summarise(group_by(bankData, default),  
           ave_balance = mean(balance),  
           num = n()  
)
```

```
## Source: local data frame [2 x 3]  
##  
##   default ave_balance   num  
## 1      no   1389.8064 44396  
## 2     yes   -137.6245   815
```



Chaining/Piping

- A `dplyr` installation includes the `magrittr` package as a dependency
- The `magrittr` package includes a pipe operator that allows you to pass the current dataset to another function
- This makes interpreting a nested sequence of operations much easier to understand



Standard Code

Code is executed inside-out.

```
arrange(filter(select(bankData, age, job, education, default), default == 'yes'), job, education, age)
```

```
## Source: local data frame [815 x 4]
```

```
##
```

##	age	job	education	default
## 1	25	admin.	secondary	yes
## 2	26	admin.	secondary	yes
## 3	26	admin.	secondary	yes
## 4	26	admin.	secondary	yes
## 5	26	admin.	secondary	yes
## 6	27	admin.	secondary	yes
## 7	27	admin.	secondary	yes
## 8	27	admin.	secondary	yes
## 9	28	admin.	secondary	yes
## 10	29	admin.	secondary	yes
##



Reformatted

```
arrange(  
  filter(  
    select(bankData, age, job, education, default),  
    default == 'yes'),  
  job, education, age)
```

```
## Source: local data frame [815 x 4]
```

```
##
```

```
##   age   job education default  
## 1  25 admin. secondary    yes  
## 2  26 admin. secondary    yes  
## 3  26 admin. secondary    yes  
## 4  26 admin. secondary    yes  
## 5  26 admin. secondary    yes  
## 6  27 admin. secondary    yes  
## 7  27 admin. secondary    yes  
## 8  27 admin. secondary    yes  
## 9  28 admin. secondary    yes  
## 10 29 admin. secondary    yes  
## .. ...      ...      ...      ...
```



With Piping

```
bankData %>%  
  select(age, job, education, default) %>%  
  filter(default == 'yes') %>%  
  arrange(job, education, age)
```

```
## Source: local data frame [815 x 4]  
##  
##   age    job education default  
## 1   25 admin. secondary    yes  
## 2   26 admin. secondary    yes  
## 3   26 admin. secondary    yes  
## 4   26 admin. secondary    yes  
## 5   26 admin. secondary    yes  
## 6   27 admin. secondary    yes  
## 7   27 admin. secondary    yes  
## 8   27 admin. secondary    yes  
## 9   28 admin. secondary    yes  
## 10  29 admin. secondary    yes  
## .. ...      ...      ...      ...
```



Pipe + group_by()

The pipe operator is very helpful for group by summaries

```
bankData %>%  
  group_by(job) %>%  
  summarise(num = n(),  
            ave_balance = mean(balance),  
            num_defaults = sum(default == 'yes'),  
            default_rate = num_defaults/num)
```

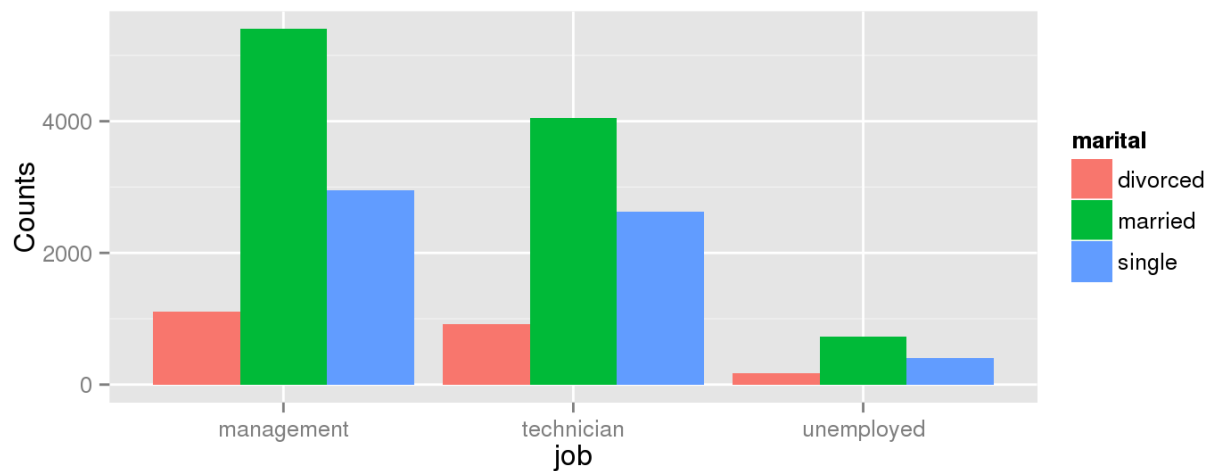
```
## Source: local data frame [12 x 5]  
##  
##       job    num ave_balance num_defaults default_rate  
## 1   admin.  5171   1135.8389         74  0.014310578  
## 2 blue-collar 9732   1078.8267        201  0.020653514  
## 3 entrepreneur 1487   1521.4701         55  0.036987223  
## 4   housemaid 1240   1392.3952         22  0.017741935  
## 5   management 9458   1763.6168        164  0.017339818  
## 6    retired  2264   1984.2151         26  0.011484099  
## 7 self-employed 1579   1647.9709         33  0.020899303  
## 8   services  4154    997.0881         75  0.018054887  
## 9    student   938   1388.0608          3  0.003198294  
## 10 technician 7597   1252.6321        130  0.017112018  
## 11 unemployed 1303   1521.7460         30  0.023023791  
## 12    unknown  288   1772.3576          2  0.006944444
```



Pipe and Plot

As a reminder, piping can also be used for non-dplyr functions.

```
library(ggplot2)
bankData %>%
  filter(job %in% c("management", "technician", "unemployed")) %>%
  group_by(job, marital) %>%
  summarise(Counts = n()) %>%
  ggplot() +
  geom_bar(aes(x = job, y = Counts, fill = marital),
           stat = 'identity', position = 'dodge')
```



Exercise

Your turn:

- Use the pipe operator to group by job and housing status
- Calculate the counts of observations, and the average and median balance



Solution

```
bankData %>%  
  group_by(job, housing) %>%  
  summarise(Counts = n(),  
            average_balance = mean(as.numeric(balance)),  
            median_balance = median(as.numeric(balance)))
```

```
## Source: local data frame [24 x 5]  
## Groups: job  
##  
##      job housing Counts average_balance median_balance  
## 1   admin.    no   1989      1378.7677         445.0  
## 2   admin.   yes   3182       983.9893         387.0  
## 3 blue-collar no   2684      1341.3308         421.5  
## 4 blue-collar yes   7048       978.8605         373.0  
## 5 entrepreneur no    618      1862.7023         390.0  
## 6 entrepreneur yes    869      1278.7986         339.0  
## 7  housemaid  no    842      1491.1876         465.5  
## 8  housemaid  yes    398      1183.3920         314.5  
## 9  management no   4780      1913.7525         600.5  
## 10 management yes   4678      1610.2076         547.0  
## ..      ...      ...      ...      ...      ...
```



Summary

mutate

Create transformations

summarise

Aggregate

group_by

Group your dataset by levels

distinct

Extract unique values (frequently used with `arrange()`)

Chaining with the `%>%` operator can result in more readable code.



Thanks for Attending!

- Any questions?
- alizaidi@microsoft.com

